FTMO_MT5 · BRIEFING FOR SIR

An AI Trading Agent for FTMO_MT5

What is already built, what the agent automates, and where you approve.

30 September 2026 · Paper trading only — live trading stays disabled

Source: FTMO_AI_Trading_Agent_Blueprint.docx — direct inspection of the FTMO_MT5 code

HOW TO READ THIS DECK

BUILT	Running in code today
REUSABLE	Written, not yet wired in
TO BUILD	Agent work still to build
YOU APPROVE	A person decides — never the agent
FIX FIRST	Unsafe today — fix before building

THE BOTTOM LINE

The engine is built. The agent is the missing operator.



5 / 10

pipeline stages already built and running in code



0

tools the agent will hold that can place an order



5

decisions that stay with a person — the agent only recommends

The agent **screens, analyses, explains and plans.** Tested code **grades, sizes and fills.** You **approve.**

STAGE BY STAGE

Where the agent works — and where you decide

1 Stock screening

TO BUILD

Ranks the FTMO-tradable universe weekly: liquidity, ATR, gap risk

You: Approve watchlist

2 Market analysis

TO BUILD

08:30 brief: earnings, macro, overnight gap exposure

No approval needed

3 Strategy selection

BUILT

Chooses only from 16 tested rules — never invents one

No approval needed

4 Backtesting

BUILT

Scheduled research; ~3,500 runs already indexed

No approval needed

5 Validation

BUILT

Explains every grade and how many tests were run

You: Promote a rule

6 Trade planning

TO BUILD

Dry-run plan before the cycle, one reason per trade

You: Approve the plan

7 Risk checks

BUILT

RiskAgent enforces FTMO limits; agent explains headroom

You: Change limits

8 Execution

BUILT

Code only. The agent holds no order tool.

You: Any live mode

9 Monitoring

TO BUILD

Watches the event log; drafts incident reports

You: Clear a HALT

10 Reporting

TO BUILD

Daily P&L, weekly memo, proposed vs executed

You: Review outbound

WHAT CHANGES DAY TO DAY

From manual steps to a daily operating rhythm

	TODAY — BY HAND	WITH THE AGENT
Screening	Tickers picked by judgement	FTMO-tradable universe ranked weekly on measured liquidity, ATR and gap risk
Market analysis	Informal reading; no written brief	08:30 brief: earnings, macro, and overnight gap exposure vs daily-loss room
Validation	Read the verdict, then decide	Every grade explained — with the number of tests behind it
Trade planning	Run the cycle, read results after	Dry-run plan before anything executes, with one reason per trade
Monitor & report	Dashboards, curl checks, reports by hand	Event-log watch, incident drafts, daily P&L and proposed-vs-executed diff

ARCHITECTURE

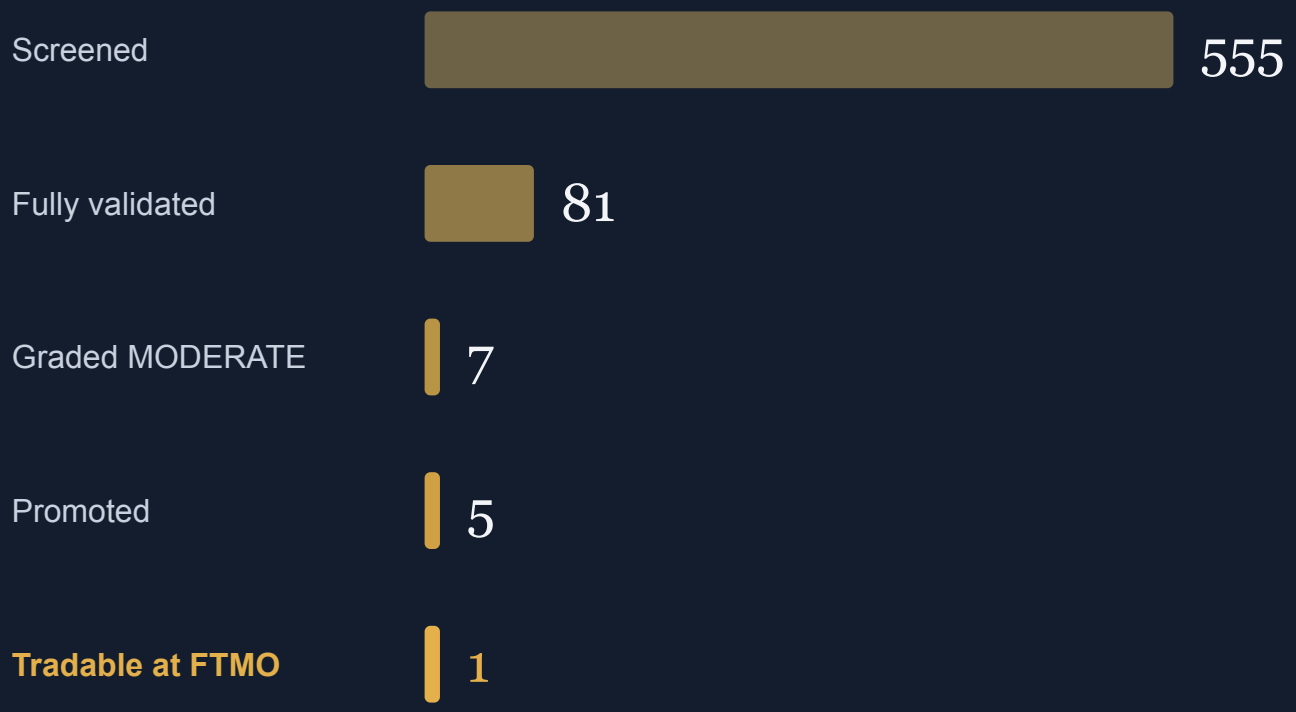
The agent proposes. Tested code decides.

YOU approval gates	**Promote a rule**	**Approve the plan**	**Clear a HALT**	**Change limits**	**Anything live**
AI AGENT own process · MCP tools only · no shell · $ cap	**Analyst** brief · calendar	**Researcher** run · read research	**Planner** dry-run · explain	**Monitor** risk state · events	**Reporter** P&L · incidents
TOOLS read or propose — never execute	**ftmo-lab** 6 tools · built	**ftmo-research** run · rank	**ftmo-paper** read · dry-run	**market data** FMP · reusable	**ftmo-reports** memo · incident
CORE built and tested — unchanged	**verdict.validate** the grade	**Systematic cycle** exits, then entries	**RiskAgent** FTMO limits · sizing	**PaperBroker** fills · ledger	**EventStore** append-only audit
BROKER	Reached only by the approved cycle. No agent tool reaches it.				

WHAT RESEARCH HAS FOUND

555 screens → 1 rule you could trade at FTMO

RULES AT EACH STAGE — LATEST RUN, 37 STOCKS × 15 RULES



WHAT IT MEANS

0 ROBUST. No rule has cleared the strict bar yet.

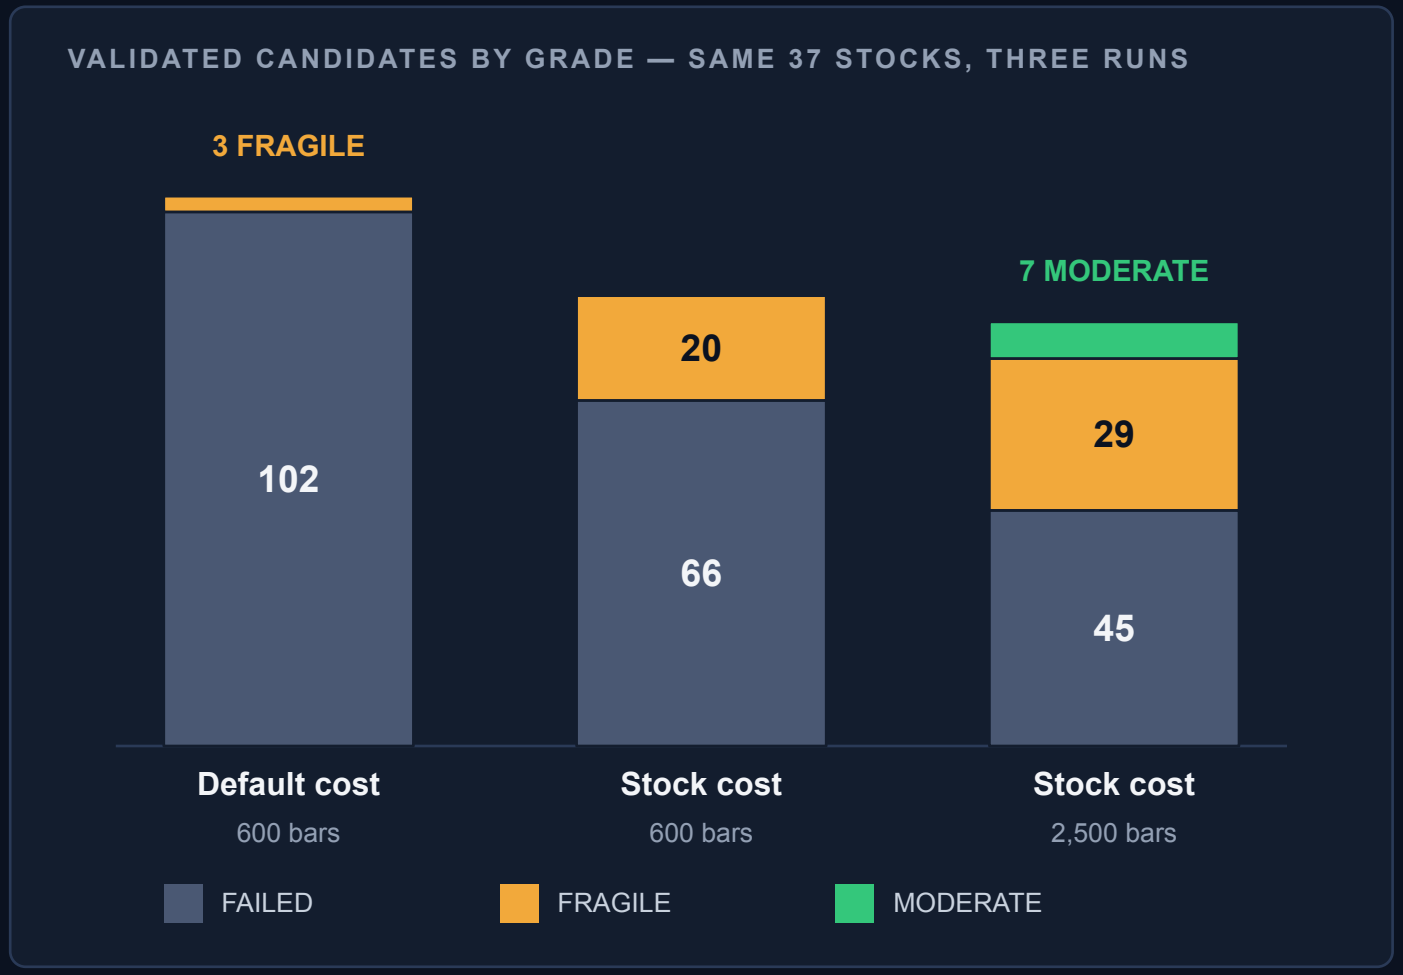
7 of 555 is not clearly better than chance — forward confirmation needed.

Only NKE survives to an FTMO-tradable rule.

Agent's first job: report this honestly, and state how many tests sit behind every pass.

RESEARCH QUALITY

Two calibration fixes changed the verdicts



70%

of gross profit eaten by a per-contract fee charged per share

$5,755 of $8,226 in a real-data replay

30 trades

needed for MODERATE — daily rules can't reach it in 600 bars

a 2,500-bar window gave the first 7

FTMO EXAMPLE

4 of the 5 promoted rules can't be traded at FTMO

VZ

Verizon

MODERATE

MEAN_REVERSION

Not listed

IEF

7–10y Treasury ETF

MODERATE

DONCHIAN_BREAKOUT

ETF — FTMO lists single stocks

TLT

20y+ Treasury ETF

MODERATE

MEAN_REVERSION

ETF — FTMO lists single stocks

GIS

General Mills

MODERATE

MEAN_REVERSION

Not listed

NKE

Nike

MODERATE

MEAN_REVERSION

Listed · Spot CFD

WHAT THE AGENT DOES

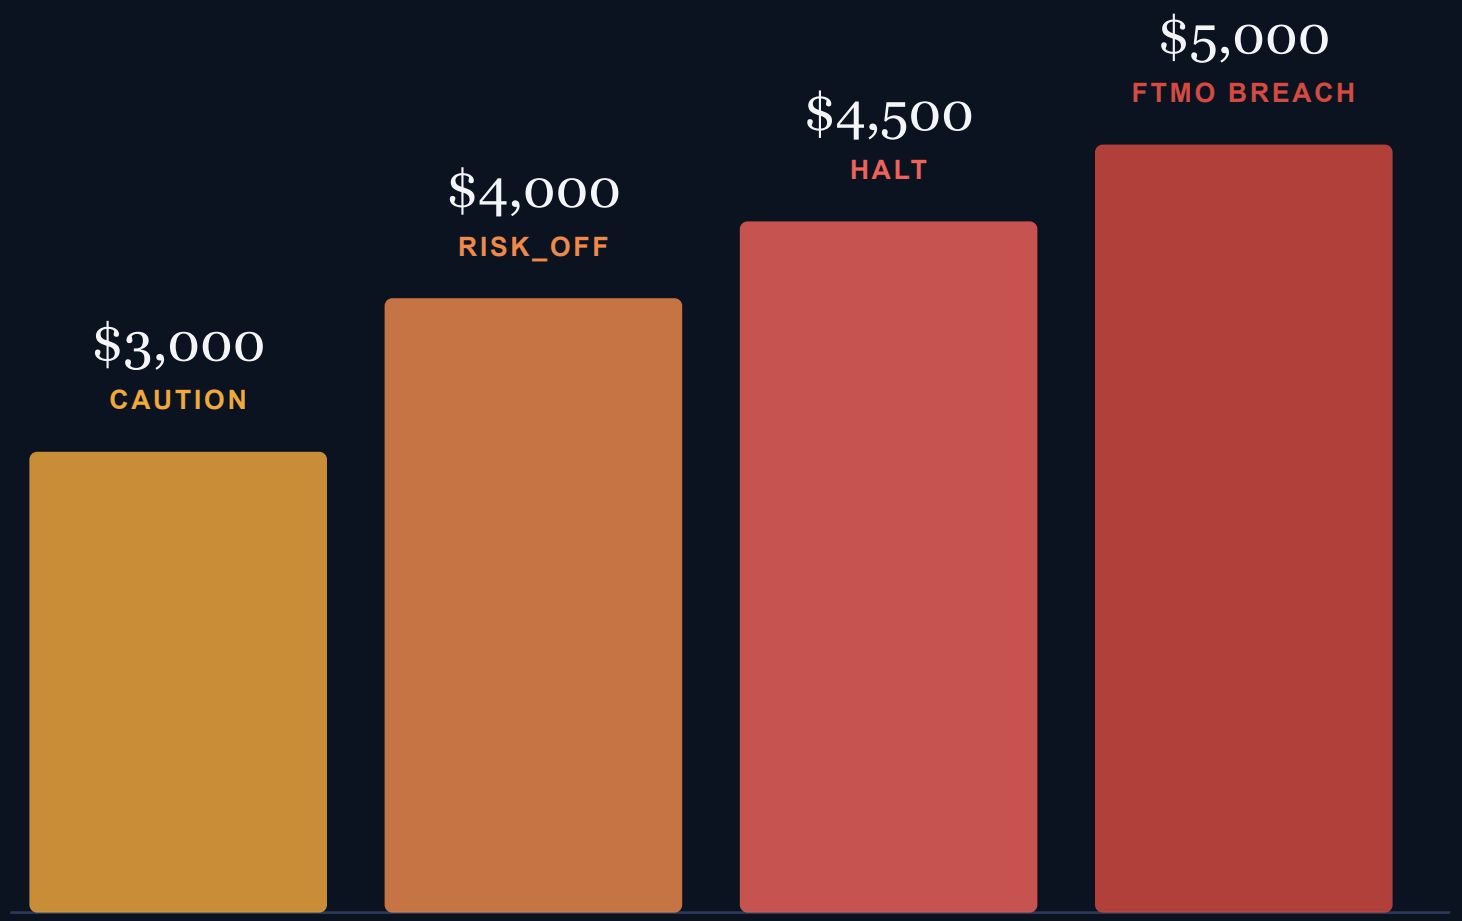
Drops the four before planning and says why in the research memo.

WHAT THE CODE SHOULD DO

Add a tradability check to the pipeline, so a grade and FTMO's list can never disagree.

FTMO RISK CONTROL

The daily-loss ladder on a $100,000 2-Step account



Daily loss that triggers each state

CODE ACTS · AGENT EXPLAINS

CAUTION

Code: New orders at half size

Agent: Alerts the team; lists open exposure

RISK_OFF

Code: Entries refused, exits allowed

Agent: Drafts incident note; proposes no entries

HALT

Code: Flatten everything, stay halted

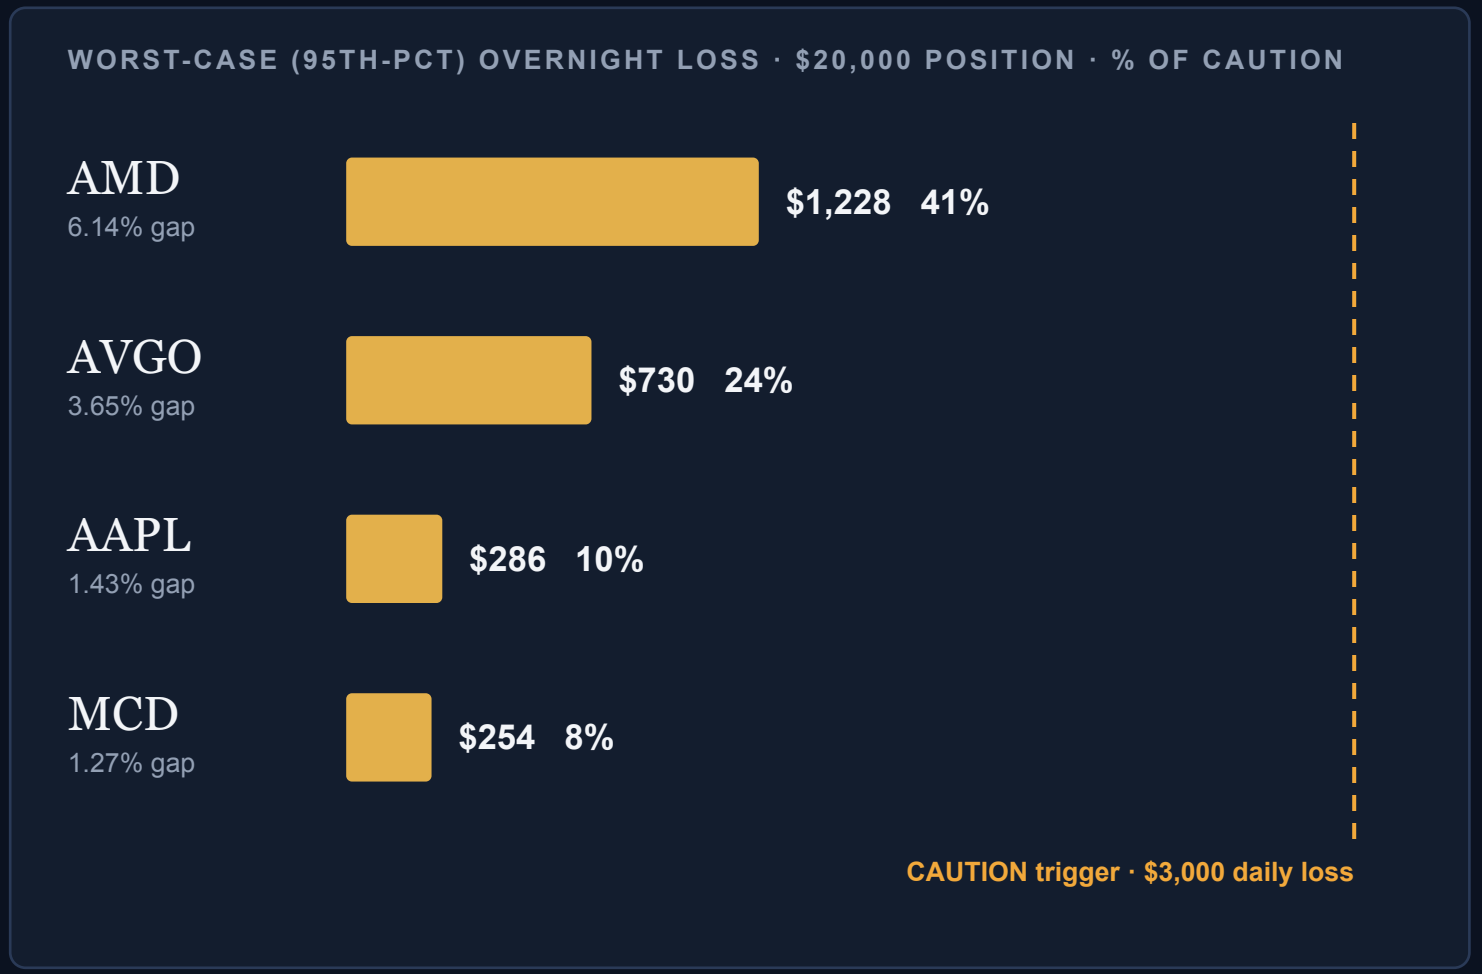
Agent: Writes the report — cannot clear it

FTMO BREACH

Code: Account fails — FTMO's own rule

Agent: —

Overnight gaps land straight on the daily-loss limit



09:35 – 16:00 ET

FTMO US share CFDs trade 16:35–23:00 GMT+3 only. A gap overnight can't be traded out of.

AGENT AT 08:30

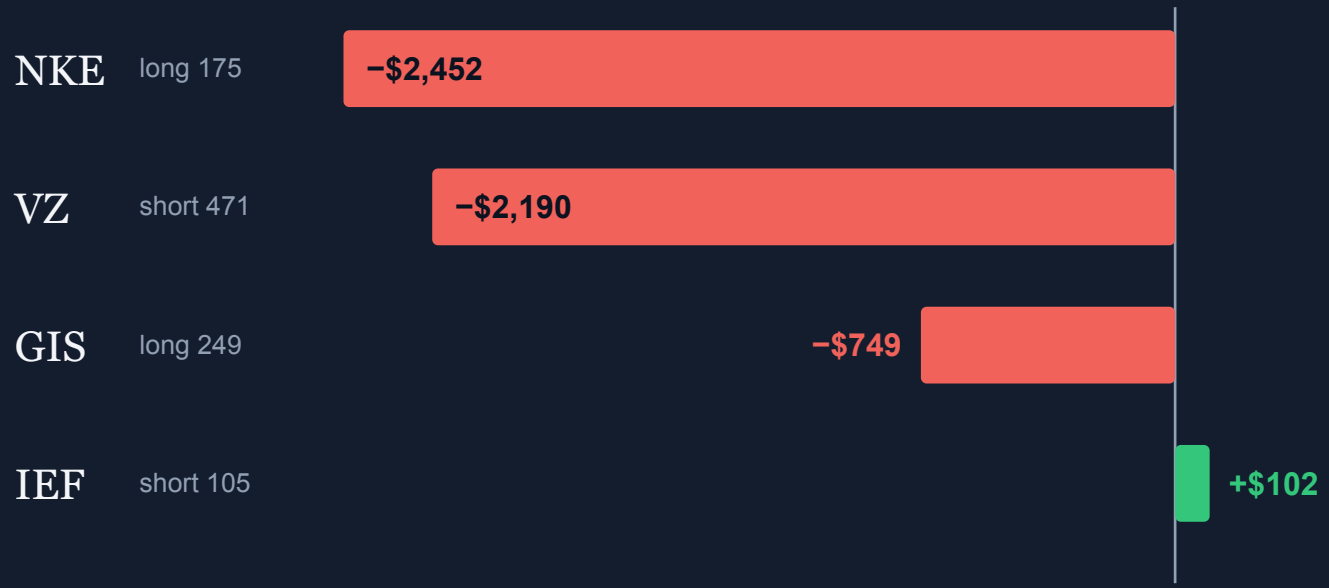
Adds up worst-case gap exposure across open positions and checks earnings. If a name reports within one session, or the total would cross CAUTION, it proposes no new entry there.

FTMO EXAMPLE

A real HALT — what the agent writes, and where it stops

$99,414	$99,693	≈ −$5,300
equity at the breach	daily-loss floor	open float, 4 positions

CLOSED P&L WHEN THE HALT FLATTENED THE BOOK



THE AGENT'S INCIDENT REPORT

CAUSE

Float from four positions counted in full against the day's allowance.

FINDING

The rules hold through drawdowns deep enough to fail a challenge.

FIX

Smaller slots or a cap on total open float — then re-run.

It cannot clear the HALT. A named person reviews and clears it.

OPERATING RHYTHM

A day with the agent (New York time)

SUNDAY	08:30	09:35 – 16:00	16:15	END OF DAY
Research + memo	**Morning brief**	**Market session**	**Plan → approve → run**	**P&L report**
Screen, validate, rank. Memo explains grades and test count.	Mode and data check, earnings, macro, gap exposure.	Agent watches the event log. Decides nothing.	Dry-run plan with reasons. RiskAgent re-checks at execution.	Realised, unrealised, risk state, proposed vs executed.
YOU PROMOTE RULES	Information only	No actions	YOU APPROVE THE PLAN	REVIEW OUTBOUND

ALWAYS ON HALT fires → RiskAgent flattens → agent writes the incident report → **a named person clears it**

CONTROL

Five decisions stay with you

Promote a rule

AGENT

Recommends, with grade and reasons

WHY YOU

Puts capital behind a rule

Approve the daily plan

AGENT

Drafts it, one reason per trade

WHY YOU

First decision on the money path

Clear a HALT

AGENT

Writes the incident report

WHY YOU

Never auto-recovers — by design

Change risk limits

AGENT

Suggests in a report

WHY YOU

Policy, not operations

Anything live

AGENT

Never

WHY YOU

Separate sign-off outside this system

Placing an order is never an agent action. The cycle executes — only after your gate, and RiskAgent re-checks at that moment.

FIX FIRST

One thing to fix before building anything

broker_hub.py · one process · 127.0.0.1:8600

Order routes

```
POST /api/{broker}/order
POST /api/{broker}/close
```

curl

/agents runner

Claude Agent SDK with shell + file edits

ACCEPTEDITS BY DEFAULT

BYPASS MODE SELECTABLE ON THE PAGE

NO TOOL ALLOWLIST

WORKING DIRECTORY SET BY THE REQUEST

Source: broker_hub.py lines 1316–1476

PHASE 0 · 2–4 DAYS

- Tool allowlist on the runner
- Remove the bypass option from the page
- Fix the working directory in code
- Replace 3 health checks that always say healthy

EXIT TEST

/agents cannot run a shell or reach an order route; each health check fails when it should.

ROADMAP

Five phases — each must pass its exit test

PHASE 0

Make it safe

Wk 0–2

Lock /agents · real health checks

EXIT TEST

Runner can't reach an order route

PHASE 1

Read-only analyst

Wk 2–5

Research, paper & report tools · daily brief

EXIT TEST

10 briefs, every number traced to a tool

PHASE 2

Proposer + you

Wk 5–9

Dry-run plans · approval in the panel

EXIT TEST

20 approved cycles; plan = execution

PHASE 3

Autonomous paper

Wk 9–15

Policy envelope · 4-week forward test

EXIT TEST

No daily-loss breach; results vs validation

PHASE 4

Live

Out of scope

Separate sign-off · capital limits

EXIT TEST

New review of every control

THE ASK

Approve Phase 0 and Phase 1

0 — **Make it safe**

Lock /agents · real health checks · 2–4 days

1 — **Read-only analyst**

Briefs and research memos — no trading · 1.5–2 weeks

WHAT THE AGENT WON'T DO

Create an edge.

Nothing grades ROBUST yet. The agent makes research faster and more honest, not profitable.

Fix position sizing.

The replay HALT came from sizing, not rule choice — that's policy work for people.

Prove profit in 4 weeks.

A short paper test catches broken plumbing, not skill.

Paper trading only — live trading stays disabled and needs a separate sign-off.